COURSE DL201 / MODULE DL

Sequence Models and Transformers

Learning goal: explain core concepts, compare related methods, and apply them to realistic machine learning or educational technology scenarios.

Document ID: DL201-04_Sequence_Models_Transformers | Format: instructional Q&A | Revision: 1.0

01. What is an RNN hidden state?

An RNN processes tokens in order and passes a hidden state between steps. That state summarizes preceding tokens, allowing sequence predictions. Long sequences are hard for basic RNNs because gradients can vanish or explode. This concept is important because it connects the underlying theory with decisions made when building a real system. Understanding both its purpose and its limitations helps avoid treating the technique as a rule that works equally well in every dataset. For example, in an educational technology system, the same idea can be tested using course content, student questions, assessment records, or retrieved passages. A team should define the expected behavior first, measure it on held-out examples, and inspect failures rather than relying only on an aggregate score. For a production or RAG-oriented workflow, preserve useful metadata and reproducible preprocessing settings so that an unexpected result can be traced back to its source. This also makes it easier to compare alternative configurations and determine whether an apparent improvement comes from better modeling, better retrieval, or simply a change in the evaluation data. An important limitation is that no single technique or metric guarantees good real-world performance. Data leakage, distribution shift, poorly chosen evaluation examples, and mismatched assumptions can all make results look stronger than they really are. The practical approach is therefore to establish a baseline, change one major factor at a time, and verify the effect using representative validation data.

02. How do LSTM and GRU handle long dependencies?

LSTMs use input, forget, and output gates plus a cell state to regulate retained information. GRUs use a simpler set of gates. Both offer learned paths for information to persist across more steps than a plain RNN. In practice, it helps to think of this as a sequence of steps rather than a single operation. Each step affects the information available to the next stage, so implementation choices should be validated on representative data. For example, in an educational technology system, the same idea can be tested using course content, student questions, assessment records, or retrieved passages. A team should define the expected behavior first, measure it on held-out examples, and inspect failures rather than relying only on an aggregate score. For a production or RAG-oriented workflow, preserve useful metadata and reproducible preprocessing settings so that an unexpected result can be traced back to its source. This also makes it easier to compare alternative configurations and determine whether an apparent improvement comes from better modeling, better retrieval, or simply a change in the evaluation data. An important limitation is that no single technique or metric guarantees good real-world performance. Data leakage, distribution shift, poorly chosen evaluation examples, and mismatched assumptions can all make results look stronger than they really are. The practical approach is therefore to establish a baseline, change one major factor at a time, and verify the effect using representative validation data.

03. What is a bidirectional LSTM?

It runs sequence processing forward and backward, combining context from both directions for each input token. It suits tagging a full sentence but cannot use future tokens for a strictly real-time autoregressive prediction. This concept is important because it connects the underlying theory with decisions made when building a real system. Understanding both its purpose and its limitations helps avoid treating the technique as a rule that works equally well in every dataset. For example, in an educational technology system, the same idea can be tested using course content, student questions, assessment records, or retrieved passages. A team should define the expected behavior first, measure it on held-out examples, and inspect failures rather than relying only on an aggregate score. For a production or RAG-oriented workflow, preserve useful metadata and reproducible preprocessing settings so that an unexpected result can be traced back to its source. This also makes it easier to compare alternative configurations and determine whether an apparent improvement comes from better modeling, better retrieval, or simply a change in the evaluation data. An important limitation is that no single technique or metric guarantees good real-world performance. Data leakage, distribution shift, poorly chosen evaluation examples, and mismatched assumptions can all make results look stronger than they really are. The practical approach is therefore to establish a baseline, change one major factor at a time, and verify the effect using representative validation data.

04. What is sequence-to-sequence modeling?

An encoder consumes an input sequence and a decoder produces an output sequence, as in translation. A fixed single-vector bottleneck can lose detail on long inputs; attention gives the decoder access to relevant encoder states. This concept is important because it connects the underlying theory with decisions made when building a real system. Understanding both its purpose and its limitations helps avoid treating the technique as a rule that works equally well in every dataset. For example, in an educational technology system, the same idea can be tested using course content, student questions, assessment records, or retrieved passages. A team should define the expected behavior first, measure it on held-out examples, and inspect failures rather than relying only on an aggregate score. For a production or RAG-oriented workflow, preserve useful metadata and reproducible preprocessing settings so that an unexpected result can be traced back to its source. This also makes it easier to compare alternative configurations and determine whether an apparent improvement comes from better modeling, better retrieval, or simply a change in the evaluation data. An important limitation is that no single technique or metric guarantees good real-world performance. Data leakage, distribution shift, poorly chosen evaluation examples, and mismatched assumptions can all make results look stronger than they really are. The practical approach is therefore to establish a baseline, change one major factor at a time, and verify the effect using representative validation data.

05. What does attention compute?

Attention forms query-key similarity scores, normalizes them into weights, and combines value vectors. The resulting representation emphasizes relevant positions. A mask can prevent attending to future tokens during causal generation. This concept is important because it connects the underlying theory with decisions made when building a real system. Understanding both its purpose and its limitations helps avoid treating the technique as a rule that works equally well in every dataset. For example, in an educational technology system, the same idea can be tested using course content, student questions, assessment records, or retrieved passages. A team should define the expected behavior first, measure it on held-out examples, and inspect failures rather than relying only on an aggregate score. For a production or RAG-oriented workflow, preserve useful metadata and reproducible preprocessing settings so that an unexpected result can be traced back to its source. This also makes it easier to compare alternative configurations and determine whether an apparent improvement comes from better modeling, better retrieval, or simply a change in the evaluation data. An important limitation is that no single technique or metric guarantees good real-world performance. Data leakage, distribution shift, poorly chosen evaluation examples, and mismatched assumptions can all make results look stronger than they really are. The practical approach is therefore to establish a baseline, change one major factor at a time, and verify the effect using representative validation data.

06. How do transformers represent order?

Self-attention combines information across token positions, while positional information tells the model where tokens occur. Encoder models read the full input; causal decoder models predict subsequent tokens using past context. In practice, it helps to think of this as a sequence of steps rather than a single operation. Each step affects the information available to the next stage, so implementation choices should be validated on representative data. For example, in an educational technology system, the same idea can be tested using course content, student questions, assessment records, or retrieved passages. A team should define the expected behavior first, measure it on held-out examples, and inspect failures rather than relying only on an aggregate score. For a production or RAG-oriented workflow, preserve useful metadata and reproducible preprocessing settings so that an unexpected result can be traced back to its source. This also makes it easier to compare alternative configurations and determine whether an apparent improvement comes from better modeling, better retrieval, or simply a change in the evaluation data. An important limitation is that no single technique or metric guarantees good real-world performance. Data leakage, distribution shift, poorly chosen evaluation examples, and mismatched assumptions can all make results look stronger than they really are. The practical approach is therefore to establish a baseline, change one major factor at a time, and verify the effect using representative validation data.