COURSE ML101 / MODULE ML

Supervised Learning

Learning goal: explain core concepts, compare related methods, and apply them to realistic machine learning or educational technology scenarios.

Document ID: ML101-02_Supervised_Learning | Format: instructional Q&A | Revision: 1.0

01. How do regression and classification differ?

Regression predicts a continuous quantity such as a score; classification predicts a category such as pass or fail. The task determines appropriate losses and metrics: MAE for score error versus precision and recall for a pass alert. A useful way to understand how do regression and classification differ is to compare the methods not only by definition, but also by the kind of input they expect, the output they produce, and the situations in which each is useful. For example, in an educational technology system, the same idea can be tested using course content, student questions, assessment records, or retrieved passages. A team should define the expected behavior first, measure it on held-out examples, and inspect failures rather than relying only on an aggregate score. For a production or RAG-oriented workflow, preserve useful metadata and reproducible preprocessing settings so that an unexpected result can be traced back to its source. This also makes it easier to compare alternative configurations and determine whether an apparent improvement comes from better modeling, better retrieval, or simply a change in the evaluation data. An important limitation is that no single technique or metric guarantees good real-world performance. Data leakage, distribution shift, poorly chosen evaluation examples, and mismatched assumptions can all make results look stronger than they really are. The practical approach is therefore to establish a baseline, change one major factor at a time, and verify the effect using representative validation data.

02. How does linear regression work?

Linear regression predicts a weighted sum of input features and an intercept. Fitting minimizes a loss such as squared error. Coefficients describe associations conditional on other features, and a straight-line relation may fail when effects are strongly nonlinear. In practice, it helps to think of this as a sequence of steps rather than a single operation. Each step affects the information available to the next stage, so implementation choices should be validated on representative data. For example, in an educational technology system, the same idea can be tested using course content, student questions, assessment records, or retrieved passages. A team should define the expected behavior first, measure it on held-out examples, and inspect failures rather than relying only on an aggregate score. For a production or RAG-oriented workflow, preserve useful metadata and reproducible preprocessing settings so that an unexpected result can be traced back to its source. This also makes it easier to compare alternative configurations and determine whether an apparent improvement comes from better modeling, better retrieval, or simply a change in the evaluation data. An important limitation is that no single technique or metric guarantees good real-world performance. Data leakage, distribution shift, poorly chosen evaluation examples, and mismatched assumptions can all make results look stronger than they really are. The practical approach is therefore to establish a baseline, change one major factor at a time, and verify the effect using representative validation data.

03. How does logistic regression produce a classification?

Logistic regression applies the sigmoid function to a weighted sum to estimate a binary-class probability. A threshold converts the probability into a label. The threshold can be adjusted when false positives and false negatives have different costs. In practice, it helps to think of this as a sequence of steps rather than a single operation. Each step affects the information available to the next stage, so implementation choices should be validated on representative data. For example, in an educational technology system, the same idea can be tested using course content, student questions, assessment records, or retrieved passages. A team should define the expected behavior first, measure it on held-out examples, and inspect failures rather than relying only on an aggregate score. For a production or RAG-oriented workflow, preserve useful metadata and reproducible preprocessing settings so that an unexpected result can be traced back to its source. This also makes it easier to compare alternative configurations and determine whether an apparent improvement comes from better modeling, better retrieval, or simply a change in the evaluation data. An important limitation is that no single technique or metric guarantees good real-world performance. Data leakage, distribution shift, poorly chosen evaluation examples, and mismatched assumptions can all make results look stronger than they really are. The practical approach is therefore to establish a baseline, change one major factor at a time, and verify the effect using representative validation data.

04. How does a decision tree choose splits?

A tree repeatedly divides examples using feature thresholds selected to reduce impurity or prediction error. Its rules are readable, but a deep tree can overfit. Maximum depth and minimum leaf size limit complexity. In practice, it helps to think of this as a sequence of steps rather than a single operation. Each step affects the information available to the next stage, so implementation choices should be validated on representative data. For example, in an educational technology system, the same idea can be tested using course content, student questions, assessment records, or retrieved passages. A team should define the expected behavior first, measure it on held-out examples, and inspect failures rather than relying only on an aggregate score. For a production or RAG-oriented workflow, preserve useful metadata and reproducible preprocessing settings so that an unexpected result can be traced back to its source. This also makes it easier to compare alternative configurations and determine whether an apparent improvement comes from better modeling, better retrieval, or simply a change in the evaluation data. An important limitation is that no single technique or metric guarantees good real-world performance. Data leakage, distribution shift, poorly chosen evaluation examples, and mismatched assumptions can all make results look stronger than they really are. The practical approach is therefore to establish a baseline, change one major factor at a time, and verify the effect using representative validation data.

05. What is a random forest?

A random forest combines trees trained on bootstrap samples, with feature subsets considered at each split. Averaging reduces the variance of a single tree. Feature importance can be informative but should be checked for bias and leakage. This concept is important because it connects the underlying theory with decisions made when building a real system. Understanding both its purpose and its limitations helps avoid treating the technique as a rule that works equally well in every dataset. For example, in an educational technology system, the same idea can be tested using course content, student questions, assessment records, or retrieved passages. A team should define the expected behavior first, measure it on held-out examples, and inspect failures rather than relying only on an aggregate score. For a production or RAG-oriented workflow, preserve useful metadata and reproducible preprocessing settings so that an unexpected result can be traced back to its source. This also makes it easier to compare alternative configurations and determine whether an apparent improvement comes from better modeling, better retrieval, or simply a change in the evaluation data. An important limitation is that no single technique or metric guarantees good real-world performance. Data leakage, distribution shift, poorly chosen evaluation examples, and mismatched assumptions can all make results look stronger than they really are. The practical approach is therefore to establish a baseline, change one major factor at a time, and verify the effect using representative validation data.

06. How do KNN and SVM differ?

K-nearest neighbors predicts from nearby training examples and needs an appropriate distance scale. A support vector machine learns a separating boundary with a margin; kernels allow nonlinear boundaries. Both need careful scaling for many feature types. A useful way to understand how do knn and svm differ is to compare the methods not only by definition, but also by the kind of input they expect, the output they produce, and the situations in which each is useful. For example, in an educational technology system, the same idea can be tested using course content, student questions, assessment records, or retrieved passages. A team should define the expected behavior first, measure it on held-out examples, and inspect failures rather than relying only on an aggregate score. For a production or RAG-oriented workflow, preserve useful metadata and reproducible preprocessing settings so that an unexpected result can be traced back to its source. This also makes it easier to compare alternative configurations and determine whether an apparent improvement comes from better modeling, better retrieval, or simply a change in the evaluation data. An important limitation is that no single technique or metric guarantees good real-world performance. Data leakage, distribution shift, poorly chosen evaluation examples, and mismatched assumptions can all make results look stronger than they really are. The practical approach is therefore to establish a baseline, change one major factor at a time, and verify the effect using representative validation data.

07. How should hyperparameters be tuned?

Choose a metric, define a search space, and evaluate candidate settings using cross validation on training data. Keep a separate test set for one final estimate. Randomized search is useful when many settings are possible and only some matter strongly. In practice, it helps to think of this as a sequence of steps rather than a single operation. Each step affects the information available to the next stage, so implementation choices should be validated on representative data. For example, in an educational technology system, the same idea can be tested using course content, student questions, assessment records, or retrieved passages. A team should define the expected behavior first, measure it on held-out examples, and inspect failures rather than relying only on an aggregate score. For a production or RAG-oriented workflow, preserve useful metadata and reproducible preprocessing settings so that an unexpected result can be traced back to its source. This also makes it easier to compare alternative configurations and determine whether an apparent improvement comes from better modeling, better retrieval, or simply a change in the evaluation data. An important limitation is that no single technique or metric guarantees good real-world performance. Data leakage, distribution shift, poorly chosen evaluation examples, and mismatched assumptions can all make results look stronger than they really are. The practical approach is therefore to establish a baseline, change one major factor at a time, and verify the effect using representative validation data.