COURSE ML101 / MODULE ML

Model Evaluation and Production

Learning goal: explain core concepts, compare related methods, and apply them to realistic machine learning or educational technology scenarios.

Document ID: ML101-04_Model_Evaluation_Production | Format: instructional Q&A | Revision: 1.0

01. How do accuracy, precision, recall, and F1 differ?

Accuracy is the share of correct predictions. Precision is TP/(TP+FP), recall is TP/(TP+FN), and F1 is their harmonic mean. For rare at-risk students, accuracy alone can hide poor recall. A useful way to understand how do accuracy, precision, recall, and f1 differ is to compare the methods not only by definition, but also by the kind of input they expect, the output they produce, and the situations in which each is useful. For example, in an educational technology system, the same idea can be tested using course content, student questions, assessment records, or retrieved passages. A team should define the expected behavior first, measure it on held-out examples, and inspect failures rather than relying only on an aggregate score. For a production or RAG-oriented workflow, preserve useful metadata and reproducible preprocessing settings so that an unexpected result can be traced back to its source. This also makes it easier to compare alternative configurations and determine whether an apparent improvement comes from better modeling, better retrieval, or simply a change in the evaluation data. An important limitation is that no single technique or metric guarantees good real-world performance. Data leakage, distribution shift, poorly chosen evaluation examples, and mismatched assumptions can all make results look stronger than they really are. The practical approach is therefore to establish a baseline, change one major factor at a time, and verify the effect using representative validation data.

02. What does a confusion matrix show?

A confusion matrix counts true positives, false positives, true negatives, and false negatives for a chosen class and threshold. It makes error types visible, so the team can judge which mistakes are more costly. This concept is important because it connects the underlying theory with decisions made when building a real system. Understanding both its purpose and its limitations helps avoid treating the technique as a rule that works equally well in every dataset. For example, in an educational technology system, the same idea can be tested using course content, student questions, assessment records, or retrieved passages. A team should define the expected behavior first, measure it on held-out examples, and inspect failures rather than relying only on an aggregate score. For a production or RAG-oriented workflow, preserve useful metadata and reproducible preprocessing settings so that an unexpected result can be traced back to its source. This also makes it easier to compare alternative configurations and determine whether an apparent improvement comes from better modeling, better retrieval, or simply a change in the evaluation data. An important limitation is that no single technique or metric guarantees good real-world performance. Data leakage, distribution shift, poorly chosen evaluation examples, and mismatched assumptions can all make results look stronger than they really are. The practical approach is therefore to establish a baseline, change one major factor at a time, and verify the effect using representative validation data.

03. What does ROC-AUC measure?

ROC-AUC summarizes how well a score ranks positives above negatives across thresholds. It does not fix a threshold or show probability calibration. For rare positives, a precision-recall curve can be more informative. This concept is important because it connects the underlying theory with decisions made when building a real system. Understanding both its purpose and its limitations helps avoid treating the technique as a rule that works equally well in every dataset. For example, in an educational technology system, the same idea can be tested using course content, student questions, assessment records, or retrieved passages. A team should define the expected behavior first, measure it on held-out examples, and inspect failures rather than relying only on an aggregate score. For a production or RAG-oriented workflow, preserve useful metadata and reproducible preprocessing settings so that an unexpected result can be traced back to its source. This also makes it easier to compare alternative configurations and determine whether an apparent improvement comes from better modeling, better retrieval, or simply a change in the evaluation data. An important limitation is that no single technique or metric guarantees good real-world performance. Data leakage, distribution shift, poorly chosen evaluation examples, and mismatched assumptions can all make results look stronger than they really are. The practical approach is therefore to establish a baseline, change one major factor at a time, and verify the effect using representative validation data.

04. When should MAE, MSE, and RMSE be used?

MAE averages absolute regression errors in target units. MSE averages squared errors and emphasizes large misses; RMSE is its square root and returns to target units. Choose based on the cost of large mistakes. The choice is context dependent. The right approach depends on the objective, the data distribution, the cost of errors, and operational constraints such as latency, memory, interpretability, and the amount of labeled data available. For example, in an educational technology system, the same idea can be tested using course content, student questions, assessment records, or retrieved passages. A team should define the expected behavior first, measure it on held-out examples, and inspect failures rather than relying only on an aggregate score. For a production or RAG-oriented workflow, preserve useful metadata and reproducible preprocessing settings so that an unexpected result can be traced back to its source. This also makes it easier to compare alternative configurations and determine whether an apparent improvement comes from better modeling, better retrieval, or simply a change in the evaluation data. An important limitation is that no single technique or metric guarantees good real-world performance. Data leakage, distribution shift, poorly chosen evaluation examples, and mismatched assumptions can all make results look stronger than they really are. The practical approach is therefore to establish a baseline, change one major factor at a time, and verify the effect using representative validation data.

05. What is cross validation?

Cross validation rotates training and validation subsets and averages performance to reduce dependence on one split. Group folds by student when records repeat; for time-series predictions, use time-ordered splits to prevent future leakage. This concept is important because it connects the underlying theory with decisions made when building a real system. Understanding both its purpose and its limitations helps avoid treating the technique as a rule that works equally well in every dataset. For example, in an educational technology system, the same idea can be tested using course content, student questions, assessment records, or retrieved passages. A team should define the expected behavior first, measure it on held-out examples, and inspect failures rather than relying only on an aggregate score. For a production or RAG-oriented workflow, preserve useful metadata and reproducible preprocessing settings so that an unexpected result can be traced back to its source. This also makes it easier to compare alternative configurations and determine whether an apparent improvement comes from better modeling, better retrieval, or simply a change in the evaluation data. An important limitation is that no single technique or metric guarantees good real-world performance. Data leakage, distribution shift, poorly chosen evaluation examples, and mismatched assumptions can all make results look stronger than they really are. The practical approach is therefore to establish a baseline, change one major factor at a time, and verify the effect using representative validation data.

06. What are data drift and model monitoring?

Data drift is a change in the distribution of incoming features; performance drift is a change in measured prediction quality. Monitor feature distributions, missingness, latency, and labeled outcomes when available. A drift alert alone does not establish that accuracy fell. This concept is important because it connects the underlying theory with decisions made when building a real system. Understanding both its purpose and its limitations helps avoid treating the technique as a rule that works equally well in every dataset. For example, in an educational technology system, the same idea can be tested using course content, student questions, assessment records, or retrieved passages. A team should define the expected behavior first, measure it on held-out examples, and inspect failures rather than relying only on an aggregate score. For a production or RAG-oriented workflow, preserve useful metadata and reproducible preprocessing settings so that an unexpected result can be traced back to its source. This also makes it easier to compare alternative configurations and determine whether an apparent improvement comes from better modeling, better retrieval, or simply a change in the evaluation data. An important limitation is that no single technique or metric guarantees good real-world performance. Data leakage, distribution shift, poorly chosen evaluation examples, and mismatched assumptions can all make results look stronger than they really are. The practical approach is therefore to establish a baseline, change one major factor at a time, and verify the effect using representative validation data.

07. What should a basic deployment include?

Package preprocessing with the fitted model, expose a versioned prediction endpoint, log request identifiers and latency, and validate input schemas. Monitor failures and output distributions; release changes with rollback capability. This concept is important because it connects the underlying theory with decisions made when building a real system. Understanding both its purpose and its limitations helps avoid treating the technique as a rule that works equally well in every dataset. For example, in an educational technology system, the same idea can be tested using course content, student questions, assessment records, or retrieved passages. A team should define the expected behavior first, measure it on held-out examples, and inspect failures rather than relying only on an aggregate score. For a production or RAG-oriented workflow, preserve useful metadata and reproducible preprocessing settings so that an unexpected result can be traced back to its source. This also makes it easier to compare alternative configurations and determine whether an apparent improvement comes from better modeling, better retrieval, or simply a change in the evaluation data. An important limitation is that no single technique or metric guarantees good real-world performance. Data leakage, distribution shift, poorly chosen evaluation examples, and mismatched assumptions can all make results look stronger than they really are. The practical approach is therefore to establish a baseline, change one major factor at a time, and verify the effect using representative validation data.